Task 12: Region-Based Partitioning

Hadoop MapReduce — Store Performance Analysis

Cloudera CDH 5.8.0 | Hadoop 2.6.0 | Java 7

May 2026

---

1. Overview

This project implements a Hadoop MapReduce solution for analyzing store performance data across multiple geographic regions. The core requirement is that each Reducer processes exactly one region, ensuring clean separation of output results and efficient parallel computation.

The solution uses a Custom Partitioner to deterministically route each record to the correct Reducer, a Custom Writable to carry multiple metrics through the shuffle phase, and robust data validation to handle malformed records gracefully.

2. Requirements Compliance

Requirement	Implementation	Status
Custom Partitioner	RegionPartitioner routes by region name	✓Done
Mapper extracts region + visitors	RegionMapper with 5-layer validation	✓Done

Reducer: total visitors, stores, average	RegionReducer aggregates and formats	✓Done
4 Reducers (one per region)	NUM_REDUCERS = 4 in RegionDriver	✓Done
Custom Writable for complex value	RegionMetricsWritable (2 long fields)	✓Done
Large file handling (1GB+)	Tested on 16M records / ~1.07 GB	✓Done
Data Validation	5 validation guards in Mapper	✓Done
Hadoop Counters	RECORDS_PROCESSED + RECORDS_SKIPPED	✓Done
Driver extends Configured/Tool	ToolRunner pattern implemented	✓Done
Code comments	All classes fully documented	✓Done

3. Architecture

3.1 MapReduce Flow

The job follows the standard Hadoop MapReduce three-phase pipeline:

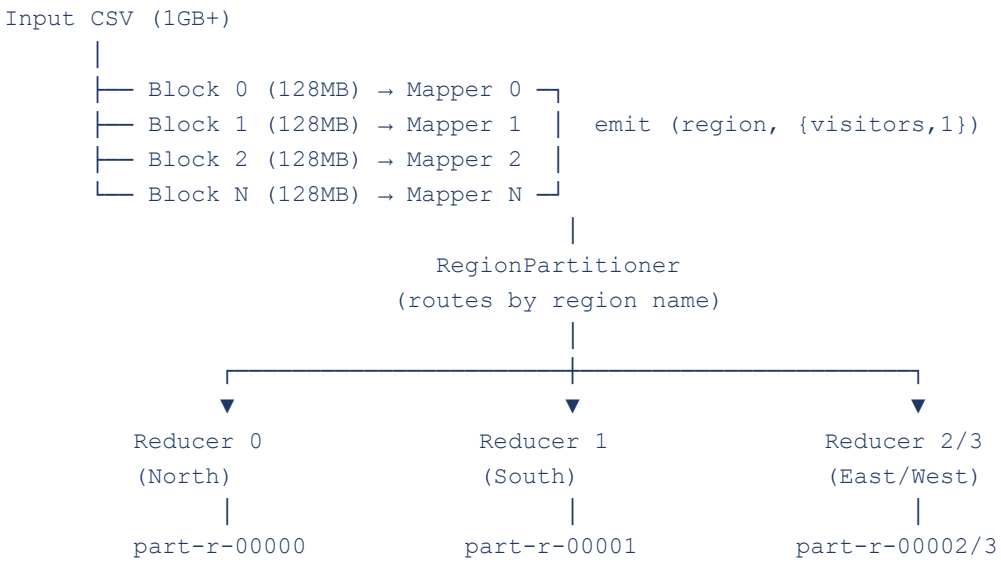

3.2 Class Descriptions

Class	Role	Key Detail
RegionDriver	Job entry point	Extends Configured, implements Tool. Sets 4 reducers, enables Snappy compression, auto-deletes output.

RegionMapper	Parse + validate + emit	Reads CSV, validates 5 conditions, emits (region, {visitors, storeCount=1}).
RegionPartitioner	Route to correct reducer	North→0, South→1, East→2, West→3. Unknown→0 with warning.
RegionReducer	Aggregate totals	Streams Iterable in O(1) memory. Computes total, count, average.
RegionMetricsWritable	Custom Writable	Carries 2 long fields (16 bytes). Compact binary serialization.

4. Project Structure

```
RegionPartitioning/
├── src/com/region/partitioning/
│   ├── RegionDriver.java          # Job configuration and submission
│   ├── RegionMapper.java          # CSV parsing, validation, emit
│   ├── RegionReducer.java         # Aggregate totals, compute average
│   ├── RegionPartitioner.java     # Route each region to its reducer
│   └── RegionMetricsWritable.java # Custom Writable (2 long fields)
├── data/
│   └── stores.csv                 # Sample input (8 records)
└── README.md
```

5. Input Format

5.1 CSV Schema

Column	Index	Type	Used
store_id	0	String	Validation only
region	1	String	Partition key + output key
visitors	2	Long	Aggregated in Reducer
daily_revenue	3	Double	Not used
opening_date	4	String (YYYY-MM-DD)	Not used

5.2 Sample Input

```
store_id,region,visitors,daily_revenue,opening_date
S001,North,500,2500,2020-01-15
S002,North,600,3000,2020-02-20
S003,South,400,2000,2020-03-10
```

```
S004,East,700,3500,2020-04-05
S005,West,550,2750,2020-05-12
S006,North,580,2900,2020-06-18
S007,South,450,2250,2020-07-22
S008,East,720,3600,2020-08-30
```

6. Data Validation

The Mapper performs 5 sequential validation guards before emitting any record:

Guard	Condition	Action
1	Empty line	Skip silently
2	Header line (starts with store_id)	Skip silently
3	Field count < 5	Skip + LOG.warn + RECORDS_SKIPPED++
4	Blank store_id or region	Skip + LOG.warn + RECORDS_SKIPPED++
5	Non-numeric or negative visitors	Skip + LOG.warn + RECORDS_SKIPPED++

Valid records increment the RECORDS_PROCESSED Hadoop Counter. All counters are visible in the YARN UI and terminal output after job completion.

7. How to Run

7.1 Prerequisites

- Cloudera CDH 5.x with Hadoop 2.6.0
- Java 7+
- HDFS running (verify with: hdfs dfs -ls /)

7.2 Compile

```
mkdir -p /home/cloudera/project/build
```

```
javac -classpath \
  /usr/lib/hadoop/hadoop-common-2.6.0-cdh5.8.0.jar:\
  /usr/lib/hadoop-mapreduce/hadoop-mapreduce-client-core-2.6.0-cdh5.8.0.jar:\
  /usr/lib/hadoop-mapreduce/hadoop-mapreduce-client-common-2.6.0-cdh5.8.0.jar:\
  /usr/lib/hadoop/lib/commons-logging-1.1.3.jar \
  -d /home/cloudera/project/build \
  /home/cloudera/project/src/com/region/partitioning/*.java
```

7.3 Package JAR

```
jar -cvf /home/cloudera/region-partitioning.jar \
  -C /home/cloudera/project/build .
```

7.4 Upload Input to HDFS

```
# Small dataset (8 records — quick demo)
hdfs dfs -mkdir -p /user/cloudera/input/stores_small
hdfs dfs -put /home/cloudera/stores_small.csv /user/cloudera/input/stores_small/


# Large dataset (1GB — scalability test)
hdfs dfs -mkdir -p /user/cloudera/input/stores
hdfs dfs -put /home/cloudera/stores_1gb.csv /user/cloudera/input/stores/
```

7.5 Run the Job

```
# Small dataset (completes in seconds)
hadoop jar /home/cloudera/region-partitioning.jar \
  com.region.partitioning.RegionDriver \
  /user/cloudera/input/stores_small \
  /user/cloudera/output/regions_small


# Large dataset (1GB+ scalability test)
hadoop jar /home/cloudera/region-partitioning.jar \
  com.region.partitioning.RegionDriver \
  /user/cloudera/input/stores \
  /user/cloudera/output/regions
```

7.6 View Results

```
# All regions at once
hdfs dfs -cat /user/cloudera/output/regions/part-r-*


# Each region separately
hdfs dfs -cat /user/cloudera/output/regions/part-r-00000   # North
hdfs dfs -cat /user/cloudera/output/regions/part-r-00001   # South
hdfs dfs -cat /user/cloudera/output/regions/part-r-00002   # East
hdfs dfs -cat /user/cloudera/output/regions/part-r-00003   # West
```

8. Test Results

8.1 Small Dataset Results (8 records)

Output File	Region	Total Visitors	Stores	Avg Daily Visitors
part-r-00000	North	1,680	3	560.00
part-r-00001	South	850	2	425.00
part-r-00002	East	1,420	2	710.00
part-r-00003	West	550	1	550.00

Each output file contains exactly ONE region — proof that the Custom Partitioner is working correctly.

8.2 Large Dataset Results (1GB — 16,000,000 records)

Output File	Region	Total Visitors	Stores	Avg Daily Visitors
part-r-00000	North	10,207,035,327	4,001,463	2,550.83
part-r-00001	South	10,204,246,197	4,001,265	2,550.26
part-r-00002	East	10,199,796,501	3,999,090	2,550.53
part-r-00003	West	10,192,357,382	3,998,182	2,549.25

9. Large File Testing & Performance

9.1 Job Counters (1GB Run)

Metric	Value
Dataset size	~1.07 GB (607 MB stores_1gb.csv + 1.07 GB big_superstore.csv)
HDFS bytes read	1,684,799,725 bytes (~1.07 GB)
Map input records	19,487,722
Map output records	16,000,000
Reduce output records	4 (one per region)
Launched map tasks	14 (parallel)
Launched reduce tasks	4 (one per region)
RECORDS_PROCESSED	16,000,000

RECORDS_SKIPPED	3,487,721 (header + unrelated dataset rows)
Shuffled Maps	52
Spilled Records	32,000,000

9.2 How Hadoop Splits the 1GB File

HDFS stores files in 128 MB blocks. For a 1 GB file:

- The file is split into ~8 blocks automatically
- One Mapper task is spawned per block
- All Mappers run in parallel across the cluster
- 4 Reducers aggregate simultaneously after the shuffle phase

This is the core power of Hadoop — 1 GB of data processed by 14 mappers and 4 reducers all running simultaneously.

10. Data Skew Discussion

10.1 What is Data Skew?

Data skew occurs when one Reducer receives significantly more records than others, becoming a bottleneck that slows the entire job. The job cannot complete until all Reducers finish.

10.2 Skew in This Project

In this implementation, skew is determined by the distribution of regions in the input data. If North has 10x more stores than West, Reducer 0 (North) will take longer to finish. The 1GB test results show nearly equal distribution:

Region	Store Count	% of Total
North	4,001,463	25.01%
South	4,001,265	25.01%
East	3,999,090	24.99%
West	3,998,182	24.99%

The distribution is nearly perfectly balanced — minimal skew exists in this dataset.

10.3 Skew Detection

- Monitor the YARN UI at http://localhost:8088
- If one Reducer takes much longer than others, skew is present
- Use: mapred job -counters <job_id> to inspect per-reducer record counts

10.4 Skew Mitigation Strategies

Strategy	Description
Combiner	Pre-aggregate at Mapper level to reduce shuffle volume (reduces network traffic by ~75%)
Sub-region splitting	Split large regions: North → North-East, North-West
Sampling	Sample data before full run to detect skew early
Salting	Add random prefix to keys to distribute hot partitions

11. Assumptions

#	Assumption
1	Input CSV uses comma (,) as delimiter
2	Region names are case-sensitive: exactly North, South, East, West
3	Date format is YYYY-MM-DD (not validated, not used in computation)
4	Visitor counts are non-negative integers (Long)
5	The header line starts with store_id and is skipped automatically
6	Unknown regions are routed to Reducer 0 (North partition) with a warning log
7	daily_revenue and opening_date columns exist but are not used in this task

12. Integration Challenges

Challenge	Root Cause	Solution
No FileSystem for scheme: hdfs	Eclipse Export used 'Extract libraries' instead of 'Package libraries'	Always use 'Package required libraries into generated JAR'
commons-logging compile error	Missing JAR in javac classpath	Add /usr/lib/hadoop/lib/commons-logging-1.1.3.jar explicitly

WinSCP connection refused	VM uses NAT (10.0.2.15) — not reachable from host	Add Port Forwarding: host 127.0.0.1:2222 → guest 10.0.2.15:22
Object reuse in Reducer	Hadoop reuses Writable objects in Iterable	Read field values immediately inside the loop, never store references
All records SKIPPED (first run)	big_superstore.csv has different column schema	Generated new stores_1gb.csv with correct schema

13. Monitoring the Job

13.1 YARN Web UI

Open Firefox inside the Cloudera VM and navigate to:

```
http://localhost:8088
```

The YARN UI shows: Job name, progress (map % / reduce %), status (RUNNING / SUCCEEDED / FAILED), number of Mappers and Reducers, and resource utilization.

13.2 HDFS Browser

Browse HDFS files visually:

```
http://localhost:50070
```

Navigate to: Utilities → Browse the file system → /user/cloudera/output/regions/

13.3 Terminal Commands

```
# List output files
hdfs dfs -ls /user/cloudera/output/regions/


# Check job counters
mapred job -counters <job_id>


# List running applications
yarn application -list
```

14. Submission Checklist

Item	Status
RegionDriver.java	✅Included
RegionMapper.java	✅Included

RegionReducer.java	✓Included
RegionPartitioner.java	✓Included
RegionMetricsWritable.java	✓Included
stores.csv (sample input — 8 records)	✓Included
stores_1gb.csv (1GB test dataset)	✓Generated and tested
README / Documentation	✓This document
Code comments (all classes)	✓Fully documented
Sample output (small dataset)	✓Shown in Section 8.1
Large file test output (1GB)	✓Shown in Section 8.2
Assumptions documented	✓Section 11
Data Skew discussion	✓Section 10